

Partitioning

Completeness	Given a fragment U derived from R, all keys in U must be in R
Disjointness	Given a fragment U derived from R, all keys in U must be unique foreign key in R
Reconstruct	$R = R_1 \cup R_2 \dots \cup R_n$
Range partitioning	$R_1 = \sigma_{A < 100}(R), R_2 = \sigma_{A \in [100, 500]}(R), R = R_1 \cup R_2$
Hash partitioning	A record is stored in node N if h(t, A) falls in N's region, adding a new node only affects on node unlike modulo method, but is oblivious to servers' capacity
Derived horizontal partitioning	$R_i = R \bowtie_A S_i$, each S_i is a partition of S for R to be complete all A in R must exist in S for R to be disjoint, S.A must be unique key for R to be disjoint and complete, R.a must be a foreign key of S with non-null values of R.a F is complete partitioning of R wrt Q, if for all fragment $R_i \in F$, either every tuple in R_i matches Q or every tuple does not match Q
Complete partitioning	F is complete partitioning of R wrt Q, if for all fragment $R_i \in F$, either every tuple in R_i matches Q or every tuple does not match Q
MTPred	$m_i = p_1^* \wedge p_2^*, MTPred(P) = m_0, m_1, \dots, m_7$ e.g. $m_0 = \neg p_1 \wedge \neg p_2 \wedge \neg p_3$ $S_i = \sigma_{m_i}(R)$, we say S_i is the min term predicate partitioning of R wrt Q and $F = \{S_1, \dots, S_m\}$, F is a complete partitioning wrt every query in Q

Query processing

RA equivalence

Commutativity of binary ops	$R \times S = S \times R, R \bowtie S = S \bowtie R$
Associativity of binary ops	$(R \bowtie S) \bowtie T = R \bowtie (S \bowtie T)$
Idempotency of unary ops	$\pi_{L'}(\pi_L(R)) = \pi_{L'}(R)$, anti-intuitive reprojection to push across join $\sigma_{p_1}(\sigma_{p_2}(R)) = \sigma_{p_1 \wedge p_2}(R)$, separate selection
Commutating sel and proj	$\pi_L(\sigma_p(R)) = \pi_L(\sigma_p(\pi_{L \cup Attr(p)}(R)))$
Push down selection	
Push down projection	

Join strategies

Strategy	Comm cost	When
Collocated	0	All partitions on the same key in all tables
Directed	size(R) if R is being repartitioned	only 1 table is not partitioned on the key
Repartitioned	size(R)+size(S)	both are not partitioned on the key
Broadcast	$(n-1) \times size(R)$ if R is broadcasted, R is smaller in size	both are not partitioned on the key

NOTE: If need to compute join from J_1 to other relation, need to see what strategy used previously. For eg. S is partitioned on B, R on C and T on F. Use broadcast join by broadcasting R for $J_1 = R \bowtie_A S$. Now, $J_1 \bowtie_B T$ only requires directed join on T since J_1 is partitioned on B already.

Storage

LSM storage

- improves write throughput by converting random I/O to sequential I/O
- append only updates
- vatch insert to disc once mem is full
- slow down read query coz need to search both mem and disk

MemTable

- updated in-pace
- flush records to disk when reaches a threshold
- SSTables are immutable and are assoc. with a range of ket values and a ts
- each new update is appended to commit log and updated to MemTable

Size-tiered compaction strategy

- SSTables in each tier have approx same size
- compaction happens when no. of SSTables reaches a threshold
- all SSTables in tier L are merged into a single SSTable in L + 1
- L becomes empty after compaction
- S_0 is more recent than $S_1, S_{0,4}$ is more recent than $S_{0,1}$

Level-tiered compaction strategy

- SSTables at L = 0 may not be disjoint
- SSTables at L > 0 have same size, disjoint
- SSTable at L i overlaps with at most F SSTables at L + 1
- compaction at L = 0, merge all SSTables with all overlapping SSTables at L + 1
- compaction at L > 0, round-robin
- NOTE: $F^{L-1} < size(L) \leq F^L$, compacts when $size(L) > F^L$. Each level stores F times as much data as previous level.

- optimize with sparse index using block's smallest key value then search with B+ tree; Bloom filter with h(x) to quickly check if a seach key exists in a SSTable block

Max size of R is mn MB. $F^{L-1} < mn \leq F^L$, using log we have $L = \lceil \log_F(mn) \rceil$. Increase F -> smaller L -> improves worst case I/O cost for searching but more I/O costs when merging SSTables.

Local indexing

- each server keeps a table to store mapping
- redundant index across servers, some sort of scatter-gather query
- easy to update since in the same server, not so good for searching

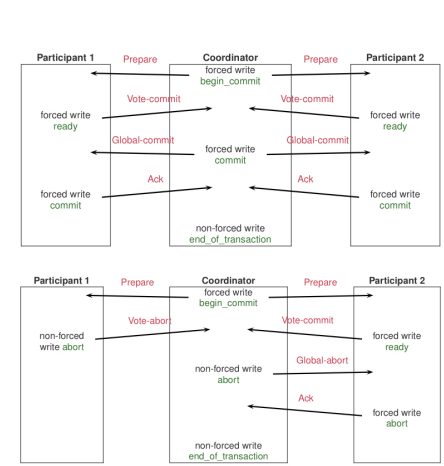
Global indexing

- every server keeps a metadata to record each index is stored at which server, partitioning done by hashing
- to search a key, go to the server with the index, get the mapping, may need to go to different servers to fetch the actual records
- good for search but update is more expensive if the mapped record is stored in different server

NOTE: To check using which index is more efficient, check no. of index partition and relation partition need to be accessed

Commit Protocol

2PC protocol



Recovery of Coordinator

Fails in state INITIAL, WAIT, ABORT
Recovery actions
Writes abort log record, sends Global-Abort to all
Does nothing if received all ACK, else sends Global-Abort to all
Does nothing if received all ACK, else sends Global-Commit

Recovery for Participants

INITIAL
READY
ABORT, COMMIT
unilaterally abort
sends Vote-Commit to TC, this is dependent
Does nothing

Basic Termination for Coordinator

timeout in state
WAIT
ABORT
COMMIT
terminate actions
Writes abort log record, sends Global-Abort to all
sends Global-Abort to non-responding participants
sends Global-Commit to non-responding participants

Basic Termination for Participants

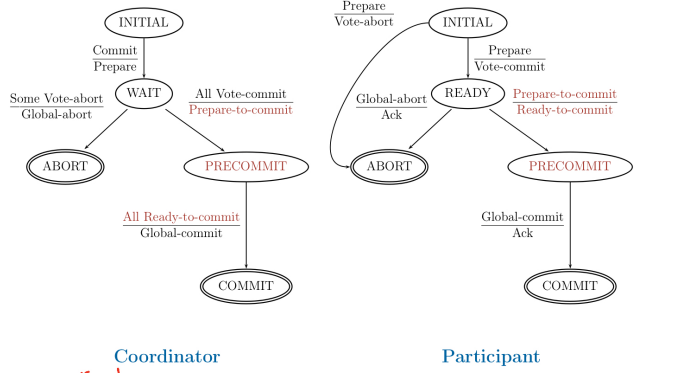
timeout in state
INITIAL
READY
terminate actions
aborts unilaterally
Blocked! Need to wait TC to recover first

Cooperative termination protocol - Participant Q receives a Decision-Request from Participant P that timeouts.

Receive in state
INITIAL
READY
COMMIT
ABORT
Replies actions
aborts unilaterally, "Aborts"
"Uncertain"
"Commit"
"Abort"

Can still be blocking if the remaining participants are all in READY state.

3PC Commit Protocol



3PC recovery protocol for TC and P

Recover in state
INITIAL
WAIT or READY, PRECOMMIT
COMMIT, ABORT
Recovery actions
aborts unilaterally
asks other TM
Does nothing
Writes abort log, sends "Global-abort"
Writes commit log, sends "Global-commit"
Sends "Global-commit" or "Global-abort" to those not yet responded

Termination protocol for P

INITIAL
READY, PRECOMMIT
3PC termination protocol 1 - Participants execute this when they timeout in READY or PRECOMMIT state
TMs elect new TC
TC sends State-Request msg to all TMs
Each TMs responds to TC about its current state
If some TMs in COMMIT, TC sends "Global-Commit" to all
else if no TMs in PRECOMMIT, TC sends "Global-Abort" to all
else (means some in READY some in PRECOMMIT), TC sends "Prepare-to-commit" to READY TMs. Sends "Global-commit" after receiving "Ready-to-commit" from these TMs.

NOTE: Failed TMs are not allowed to join Termination protocol to prevent inconsistent global decision. 3PC1 is correct if there is no connection failure. 3PC1 is blocking but correct in the event of total site failure, need to wait for a TM to recover such that it can recover independently(in INITIAL, ABORT, COMMIT) or it was the last TM to fail. The former would notify using its recovered state, the latter would execute termination protocol. In the event of comm failure, execute 3PC2.

3-PC termination protocol 2 - to ensure correctness when comm failure

- TMs elect new TC
- TC sends State-Request msg to all TMs
- Each TMs responds to TC about its current state
- If some TMs in COMMIT, TC sends "Global-Commit" to all
- else if some in ABORT, TC sends "Global-Abort" to all
- else if majority in READY/PRECOMMIT, same as 3PC1. If PRECOMMIT and Ready-to-commit forms majority, sends "Global-commit". Otherwise TC and TMs are blocked.
- else if majority in READY/PRECOMMIT/PREABORT, sends "Prepare-to-abort. If PREABORT and "Prepare-to-abort" forms majority, sends "Global-abort". Otherwise, TC and TMs are blocked.
- else TC and TMs are blocked

NOTE: 3PC2 is correct when total site failure and comm failure, is non-blocking in the absence of comm failure as long as majority of TMs are still operational

Concurrency Control

View equivalent	Same read from and final Write
Conflict equivalent	Same ordering of conflicting pairs
View/Conflict serializable	S is view/conflict equivalent to a serial schedule
Serializable global schedule	Each local schedule is view/conflict serializable -> exists a local serial schedule Each local serial schedule is a subsequence of an overall global serial schedule -> each local serial schedule is compatible ordering wise
monoversion schedule	each read action returns the most recently created object version
multiversion view equivalent	Read from only, ignore final write. $W_i(x)$ creates a new version of x denoted by x_i, x_0 is initial version
serial multiversion schedule	a multiversion schedule that is multiversion view equivalent to a serial monoversion schedule
recoverable schedule	if T reads from T' , T must commit after T' for S to be recoverable

NOTE MVSS is not necessarily VSS, but VSS must be MVSS

2PL Lock based CC protocol

- To read O, a txn must holde a S-lock or X-lock on O
- To write to O, a txn must hold X-lock on O
- A txn is not allowed to request for any more locks once it releases a lock, thus 2PL is guaranteed conflict serializable

- Strict 2PL says a txn must hold on to locks until it commits or aborts, thus S2PL is guaranteed recoverable
- Deadlock is detected using wait-for graph or timeout mechanism.
- Wait-die policy: non-preemptive, txn gets aborted only when requesting for a lock, a txn with all the locks it needs is never aborted and younger txns may get aborted repeatedly
- wound-wait: preemptive
- lock-conversion to allow for interleave schedule, but need to check before granting conversion

Multiversion based CC protocol

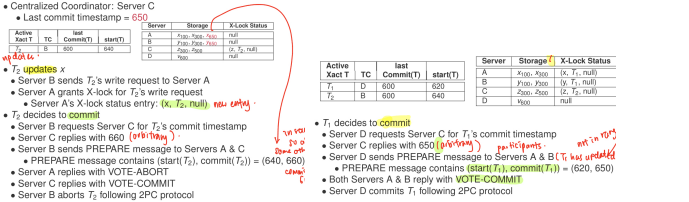
- Benefits: Read-only txns never get aborted and are not blocked by update txns, update txns are never blocked by read-only txns
- Concurrent txns: $[start(T), commit(T)] \cap [start(T'), commit(T')] \neq \emptyset$
- Only allow one txn to commit if multiple concurrent txns updated the same object to ensure serializability. (Disjoint write-sets)
- First-commit win vs First-update win. In FUW if x-lock is not held by T' , txn is granted X-lock. Txn aborts if O has been updated by concurrent txns, otherwise txn proceeds. If X-lock is held by T' and it has aborted, txn is granted X-lock and same as previous. If T' commits, txn is aborted. Txn releases X-lock after commit or abort.
- SI does not guarantee MVSS, must use serializable SI for that

Centralized vs Distributed 2PL

Lock management	CC even if it doesn't store the object	D2PL	Managed by each site's lock manager
Comm cost	ReqLock and RelLock to CC, read req to site storing object, read res to coord TM	Read req to site A storing O, site A grants S-lock and Read res to coord TM, coord TM sends RelLock to site A	central deadlock detector to construct global WFG from local WFG sent by each site
Deadlock	Aware of all deadlocks		

Centralized Snapshot Isolation

- One site as CC, say site A stores x
- Start a new txn T: TC req start(T) and lastCommit(T)
- $R_i(x)$: TC sends lastCommit(T) to TM_A , TM_A responds with latest version wrt lastCommit(T)
- $W_i(x)$: TC sends write req to TM_A , TM_A checks if X-lock can be granted and responds accordingly, T is blocked if cannot get lock
- Commit(T): TC sends req to CC for commit(T). TC execs 2PL with start(T) and commit(T). When P receives "PREPARE", it checks for WW-conflicts btwn T and committed concurrent Txns, responds "Vote-to-abort" if found concurrent version



Replication

Mutual consistency and 1-copy serializability

- S_{RD} is equivalent to S_{1C} if
 - T_j reads x from T_i in S_{RD} iff T_j also reads x from T_i in S_{1C} and
 - for each final write $W_i(x)$ in S_{1C} , $W_i(x)$ is a final write in S_{RD} for some copy x_A of x. MC if all copies of x.

Eager Strong MC	Lazy
Updates other replicas asynchronously using refresh txn	Txn commits as soon as one replica is updated
Updates from different txns can conflict	Updates from different txns can conflict
Need to ensure updates are applied in same order	Need to ensure updates are applied in same order

For all protocols, we assume strict 2PL and statement propagation

Centralized each primary or master copy site runs a lock manager for S2PL

Update is applied to master copy first before propagating to slave copies

If need latest data, must read master copy

Master/Primary site is bottleneck

Write must be done on master site first, read can be done locally

S2PC runs among involved primary sites

S2PC runs among sites that were updated

To determine if schedules could be produced by which protocol, do the following steps:

1. Check for refresh txn, if yes then lazy
2. Check for 1-copy serializability, if no then cannot be eager
3. Check where txn is issued. Then check where reads and writes are performed for centralized or distributed

Different operations using different protocols:

1. $R_i(x)$: Eager Centralized and Lazy Centralized (Assuming T_i is issued at S_B , primary copy at S_A)

- TM_B requests S-lock from TM_A . Blocked if not granted
 - Read locally if local copy exists, else sends read to other site
- Eager Distributed and Lazy Distributed
- If local copy exists, request S-lock locally. Blocked if S-lock not granted. Read locally
 - Send read to other site, say S_A . If S-lock granted, value returned from S_A
2. $W_i(x)$:
- Eager Centralized and Lazy Cnetralized
- TM_B requests X-lock from TM_A . Blocked if not granted
 - TM_B executes $W_i(x)$ and notifies TM_B lock is granted
 - If Eager, TM_B executes $W_i(x)$ and sends $W_i(x)$ to TM_C . Ignore this step for lazy protocol
- Eager Distributed
- if has local copy, requests for X-lock locally. Blocked if not granted.
 - TM_B executes $W_i(x)$ locally and sends $W_i(x)$ to TM_A and TM_C .
 - TM_A and TM_C grants X-lock and executes $W_i(x)$. Blocked if not granted.
 - if no local copy, sends $W_i(x)$ to site with x, say S_A .
 - TM_A grants X-lock to execute $W_i(x)$. TM_A notifies TM_B and sends $W_i(x)$ to TM_C
- Lazy Distributed
- if has local copy, requests for X-lock locally and executes $W_i(x)$ locally.
 - Blocked if not granted.
 - if no local copy, sends $W_i(x)$ to site with x, say S_A .
 - TM_A grants X-lock to execute $W_i(x)$.
3. $Commit_i(T)$:
- Eager Centralized and Eager distributed run 2PC among the servers involved.
- Lazy Centralized:
- TM_B sends commit(T_i) to TM_A .
 - TM_A executes commit(T_i) and releases locks for T_i .
 - TM_A checks if X-locks can be granted for T_i 's refresh txns
 - if granted, TM_A sends refresh txns to other sites to propagate T_i 's updates.
 - otherwise propagation of refresh txns is blocked
 - Slave sites must ensure refresh txns are applied in the same order using commitTS. This is trivial since TM_A should know the commitTS
- Lazy Distributed:
- if TM_B executed $W_i(x)$ locally, TM_B will execute commit
 - TM_B releases X-lock for T_i , and sends T_i 's $W_i(x)$ to TM_A and TM_C
 - TM_A and TM_C each grants X-lock for $W_i(x)$ and executes $W_i(x)$.
 - Important to have reconciliation of inconsistent updates.
 - Use Las-Write heuristics (aka timestamp order heuristic) to apply updates. Only works for blind-writes
 - If 2 concurrent updates $W_i(x)$ and $W_j(x)$, $W_i(x)$ wins if $TS(T_i) < TS(T_j)$
 - $W_j(x)$ will be ignored by TM. This will not work if $W_j(x)$ is not a blind write.

Handling failure

- Detect site failures with timeout mechanism
- If slave site fail:
 - Lazy replication: Sync unavailable replicas later when they are available
 - Eager replication: ROWA cannot terminate if one replica is down. Use ROWAA (ROWA Available) to only update available replicas and terminate txn -> sync unavailable replicas when they are available
- If master site fail:
 - Wait for master to recover -> bad availability
 - Elect new master -> might have split brain problems in the event of network partition
 - CAP theorem: When there is network partition, forfeit consistency or availability, can't have both.

Quorum consensus protocol

- $T_r(O) + T_w(O) > Wt(O)$
- $2 \times T_w(O) > Wt(O)$
- To read an object O, acquire S-lock on a read quorum $Q_r(O)$ where $Q_r(O) \geq T_r(O)$, read all copies in Q and return highest version copy
- To write an object O, acquire X-lock on a write quorum $Q_w(O)$ where $Q_w(O) \geq T_w(O)$, write all copies in Q and update their version to n + 1
- For a system with n servers to be k-tolerant where k server has failed, w is the weight of each server:
 - $(n - k)w \geq \max(T_w(O), T_r(O))$ for the system to be functional
 - Solve for n to know the min number of servers for k-tolerant

Consistency level

In Pileus, we use 2PC commit protocol and distributed SI protocol concurrency control To determine readTS of a txn:

- Client selects a MART based on the desired consistency level
 - Monotonic read: Max commitTS of All previous Get regardless of object in current sesh
 - Read-my-writes: Max commitTS of previous put of object in current sesh
 - Causal: Max commitTS of put and get regardless of objects in current sesh
 - Eventual: 0
 - Bounded staleness: $\text{realTimeToLogicalTime}(\text{client's clock time} - \text{realTime})$
 - Strong: Max CommitTS of latest version of all objects in all sesh
- Get servers that has highTS geq MART(T) of all objects
- Each object, select the server with lowest latency, if tie then select the server with larger highTS
- readTS = min of all highTS

End txn commit protocol

- client selects a commit coordinator among participants, where participants are pServers with data updated by txn
- client sends commit req to CC using readTS(txn), Put-set (PS), Largest commit timestamp (LCT) among all PS
- CC updates its localClock = max(localClock, LCT+1), and sends "Prepare-to-commit" with PS_i to P_i
- P responds with proposedTS and appends (proposedTS, PS) to pending list
- CC picks largest proposedTS as commitTS(txn), and sends commitTS to all P.
- P checks commitTS against SI protocol for concurrent updates, responds "Vote-to-abort" or "Vote-to-commit"

- CC force-write log before sending "Global-commit" and all PS, or non-force-write for "Global-abort. CC informs client that txn has committed.
- P_i processes PS_i by creating new version of object upon receiving "Global-commit" and appends to propagate list. Sends "ACK" to CC, removes txn's entry from pending list, asynchronously propagates refresh txn to secondary servers from propprage list.

Operations blocked by pending txn T' in pServer:

- Get(k) by T if T' has updated k and T' 's proposedTS \leq readTS(T) -> commitTS(T') leq readTS(T)
- both have updated the same key and T' 's proposedTS \leq commitTS(T) -> commitTS(T') leq commitTS(T)
- Replicating updates of T for propagation, same reasoning as above.

Raft Consensus algorithm

- 2PC and 2PL use RSM to replicate log files. i.e. when a coordinator crashes in 2PC, we can use RSM to replicate log file to other replicas for fault tolerance. Centralized lock manager maintains a lock table for 2PL, can be replicated for fault tolerance and prevent blocking

Terms

- each term starts with election
- at most one leader can be elected per term
- each server maintains **currentTerm**, **votedFor**, **log[(index,term,command)]**
- server replies RPC with its term and updates its current term number upon receiving larger term number
- Election safety and liveness: each follower can only vote once per term. Some leader must eventually be elected, thus election timeout >> broadcast time

RPC and Timer

- More complete log: X is more complete than Y
 - X.lastLogTerm > Y.lastLogTerm OR
 - (X.lastLogTerm = Y.lastLogTerm) AND (X.lastLogIdx > Y.lastLogIdx)
- RequestVote(candidateId,term,lastLogIdx,lastLogTerm) (term,voteGranted)
- Election timer, Leader Timer, Client Timer: If latency is expected to increase for i.e. geographical latency, need to increase timer as well
- RequestVote is success iff (P = R's log is not more complete than sender's):
 - RPC.term > R.currentTerm AND P
 - RPC.term = R.currentTerm AND R.votedFor = null AND P
 - R.votedFor = RPC.candidateId AND RPC.term = R.currentTerm
- NOTE: If RPC.term > R.currentTerm but $\neg P$, then R.currentTerm = RPC.term but R.votedFor = null. These rules ensure leader's log is at least as complete as a majority of servers'
- Log matching property: 2 entries with same index and term:
 - store the same command
 - preceding entries are identical
- Leader only executes committed log. Log is deemed committed if the **leader that created it** has replicated the log to majority of servers
- Volatile states of all servers: commitIdx, lastAppliedIdx; All leaders: nextIdx[], matchIdx[]
- AppendEntries(leaderId,leaderTerm,leaderCommit,prevLogIdx,prevLogTerm,entries[]) (term,succes)
- F processes AppendEntries if not a heartbeat msg:
 - set F.currentTerm = leaderTerm if leaderTerm > currentTerm
 - if F's does not contain prevLogIdx and prevLogTerm, reply (currentTerm,false) so leader can back off and send older logs over
 - else, delete conflicting entries and append logs
 - set F.commitIdx = min(leaderCommit, F.lastLogIdx) if leaderCommit > commitIdx
- Rules for all servers:
 - commitIdx > lastApplied, increment lastApplied by 1 and apply log[lastApplied] to RSM
 - step down to follower if received term > currentTerm
- Rules for leaders:
 - Only respond to client after commitment and execution of command
 - If AppendEntries is successful, update nextIdx and matchIdx. Otherwise decrement nextIdx and resend older entries
 - set commitIdx = N if logs are replicated to majority of server, such that N > commitIdx, matchIdx[i] \geq N and log[N].term = currentTerm
- Election safety, Election liveness, Leader Append-Only (never deletes its own logs), Log matching, Leader Completeness, State machine safety(once executed, no other servers will apply different log entry at the same index)

Query optimization

$$SF(\sigma_p(R)) = \frac{\text{card}(\sigma_p(R))}{\text{card}(R)} \quad SF(R \bowtie S) = \frac{\text{card}(R \bowtie S)}{\text{card}(R) \times \text{card}(S)}$$

$$SF(R \ltimes S) = \frac{\text{card}(R \ltimes S)}{\text{card}(R)}$$

Join selectivity estimation $SF(R \ltimes A \ S) \approx \frac{1}{\max(\text{card}(\pi_A(R)), \text{card}(\pi_A(S)))}$

$$SF(R \ltimes A \ S) \approx SF_{SJ}(S.A) = \frac{\text{card}(\pi_A(S))}{|\text{domain}(A)|}$$

Optimizing total cost

$R \ltimes A \ S$ removes dangling tuple in R wrt to $R \bowtie S$. Beneficial semijoin happens if benefit > cost where,

$$\text{Cost}(R \ltimes A \ S) = T_{\text{msg}} + T_{tr} \times \text{size}(\pi_A(S)) \text{ to send } \pi_A(S) \text{ over to site R}$$

$$\text{Benefit}(R \ltimes A \ S) = T_{tr} \times \text{size}(R) \times (1 - SF_{SJ}(S.A))$$

Comparable plans

1. both plans have the same output schema
 2. final operator for both plans is executed at the same site3.
 3. both plans output are either unordered or sorted in the same order
- Basic vs Enhanced DP**
1. basic does not allow for semijoin
 2. enhanced requires more processing to handle semijoin:
 - Enumerator extension since left and right operand might not be disjoint
 - Avoid redundant join and semijoin (every leaf to root path appears once only)
 - Fix-point iteration to make plans complete
 - Vertical pruning such as inter and intra-entry pruning, as comparable plans may be stored across different optPlan() entries